

WEKA ASSIGNMENT

BY: AARON PULIKKOTIL

1. DATASET DESCRIPTION:

Dataset Name: Labor

Source of the dataset:

<https://github.com/tertiarycourses/Weka/blob/master/Weka%20datasets/labor.arff>

Dataset description: The data set contains 1 classes of 57 instances each where each class refers to a type of iris plant.

No of attributes: 17

No of instances: 57

2. INPUT & OUTPUT:

LABOUR	1.Duration	Duration: class
	2.Wage increase 1 st year	Wage increase 1 st year: class
	3.Working hours	Working hours: class
	4.Pension	Pension: class
	5.Education allowance	Education allowance: class
	6.Statutory Holidays	Statutory Holidays: class
	7.Vacation	Vacation: class
	8.Standby Pay	Standby Pay: class
	9.Bereavement Assistance	Bereavement Assistance: class
	10.Class	

3. DATA PREPARATION:

Removal of null data types, empty instances and false attribute.

4. SELECTING AN ALGORITHM:

Algorithm	Result from Weka	
	Correctly Classified Instances	Mean Absolute Error
Weka Classifiers		
NaiveBayes	51	0.1042
NaiveBayesMultinomialText	37	0.4574
NaiveBayesUpdateable	51	0.1042
Weka Classifiers Functions		
Logistic	53	0.0641
MultilayerPerceptron	49	0.1524
Weka Classifiers Meta		
AdaBoostM1	50	0.1519
Bagging	48	0.2979
Weka Classifier Trees		
J48	42	0.3192
RandomTree	45	0.2438
Weka Classifier Rules		
DecisionTable	43	0.2996
JRip	44	0.2256

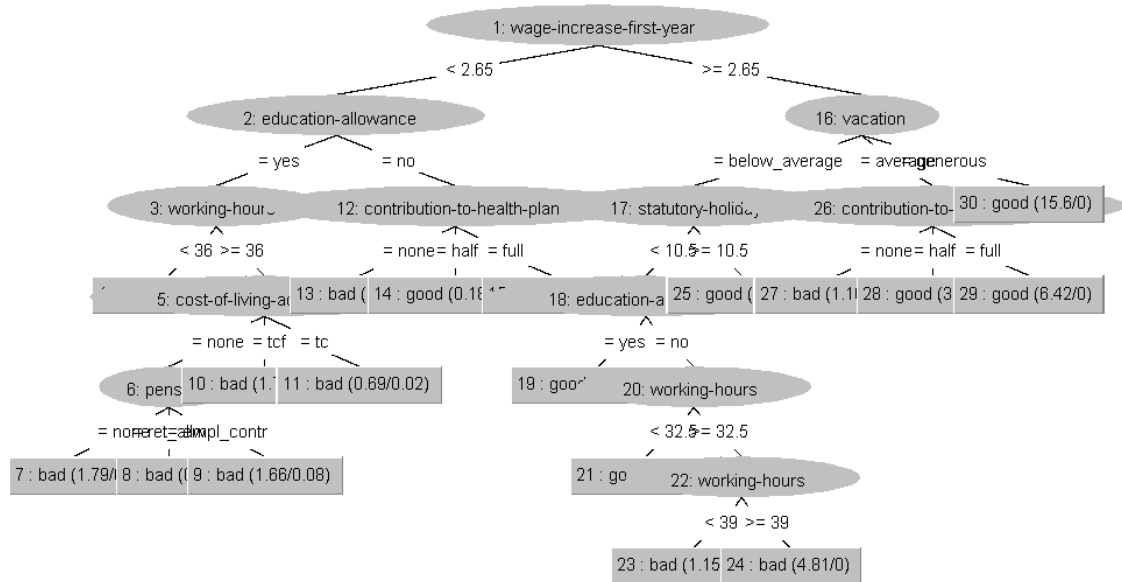
5.DATA PREPROCESSING:

Attribute Evaluator	Classifier Subset Eval (Classifier: meta.Bagging with training parameter: classifiers = J48 and the others in their default)
Search Method	BestFirst (Direction: Forward)
Evaluation Mode	10-fold cross-validation (Folds = 10, Seed =1)

6.TREE PRODUCED:

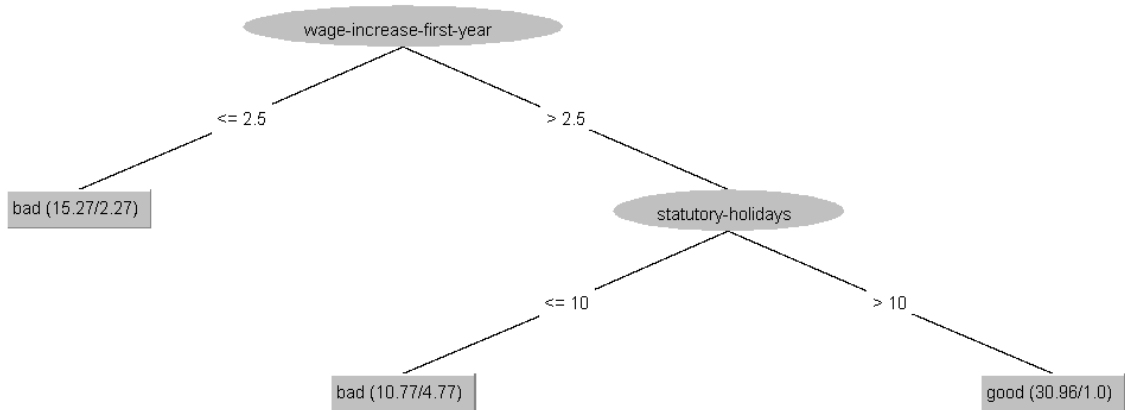
Weka Classifier Tree Visualizer: 16:30:08 - trees.RandomTree (labor-neg-data)

Tree View



Weka Classifier Tree Visualizer: 16:30:44 - trees.J48 (labor-neg-data)

Tree View



7. SCREENSHOTS:

The screenshot shows the Weka Explorer interface with the 'Classify' tab selected. The classifier is 'JRip - F 3 -N 2.0 -O 2 -S 1'. The test options are set to 'Cross-validation' with 'Folds' set to 10. The result list on the left shows '135149 - bayes.NaiveBayes' selected. The classifier output displays the following information:

Time taken to build model: 0.01 seconds

==== stratified cross-validation ====

==== Summary ====

Metric	Value
Correctly Classified Instances	51
Incorrectly Classified Instances	6
Kappa statistic	0.7741
Mean absolute error	0.1042
Root mean squared error	0.2637
Relative absolute error	22.7763 %
Root relative squared error	55.2266 %
Total Number of Instances	57

==== Detailed Accuracy By Class ====

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
0.900	0.100	0.018	0.900	0.857	0.776	0.965	0.926	bad	
0.892	0.100	0.543	0.892	0.517	0.776	0.965	0.903	good	
Weighted Avg.	0.895	0.103	0.899	0.895	0.896	0.776	0.965	0.963	

==== Confusion Matrix ====

a \ b	bad	good
18 2 a = bad	18	2
4 33 b = good	4	33

The screenshot shows the Weka Explorer interface with the 'Classify' tab selected. The classifier is 'JRip - F 3 -N 2.0 -O 2 -S 1'. The test options are set to 'Cross-validation' with 'Folds' set to 10. The result list on the left shows '135220 - bayes.NaiveBayesMultinomialText' selected. The classifier output displays the following information:

Time taken to build model: 0 seconds

==== stratified cross-validation ====

==== Summary ====

Metric	Value
Correctly Classified Instances	37
Incorrectly Classified Instances	20
Kappa statistic	0
Mean absolute error	0.4574
Root mean squared error	0.4775
Relative absolute error	100 %
Root relative squared error	100 %
Total Number of Instances	57

==== Detailed Accuracy By Class ====

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
0.000	0.000	?	0.000	?	?	?	0.472	0.339	bad
1.000	1.000	0.649	1.000	0.787	?	?	0.472	0.637	good
Weighted Avg.	0.649	0.649	?	0.649	?	?	0.472	0.532	

==== Confusion Matrix ====

a \ b	bad	good
0 20 a = bad	0	20
0 37 b = good	0	37

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose JRip - F 3 -N 2.0 -O 2 -S 1

Test options

☐ Use training set

☐ Supplied test set

☒ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Nom) class

Start Stop

Result list (right-click for options)

135149 - bayes.NaiveBayes

135220 - bayes.NaiveBayesMultinomialText

135425 - bayes.NaiveBayesUpdateable

162623 - functions.Logistic

162645 - functions.MultilayerPerceptron

162746 - meta.AdaBoostM1

162808 - meta.Bagging

163008 - trees.RandomTree

163044 - trees.J48

163228 - rules.DecisionTable

163249 - rules.JRip

Classifier output

Time taken to build model: 0 seconds

==== Stratified cross-validation ====

==== Summary ====

	Correctly Classified Instances	51	89.4737 %
	Incorrectly Classified Instances	6	10.5263 %
	Kappa statistic	0.7741	
	Mean absolute error	0.1042	
	Root mean squared error	0.2637	
	Relative absolute error	22.7763 %	
	Root relative squared error	55.2266 %	
	Total Number of Instances	57	

==== Detailed Accuracy By Class ====

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.900	0.108	0.818	0.900	0.857	0.776	0.965	0.926	bad
	0.892	0.100	0.943	0.892	0.917	0.776	0.965	0.983	good
Weighted Avg.	0.895	0.103	0.899	0.895	0.896	0.776	0.965	0.963	

==== Confusion Matrix ====

	a	b	<-- classified as
18	2	1	a = bad
4	33	1	b = good

Status

OK

Log

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose JRip - F 3 -N 2.0 -O 2 -S 1

Test options

☐ Use training set

☐ Supplied test set

☒ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Nom) class

Start Stop

Result list (right-click for options)

135149 - bayes.NaiveBayes

135220 - bayes.NaiveBayesMultinomialText

135425 - bayes.NaiveBayesUpdateable

162623 - functions.Logistic

162645 - functions.MultilayerPerceptron

162746 - meta.AdaBoostM1

162808 - meta.Bagging

163008 - trees.RandomTree

163044 - trees.J48

163228 - rules.DecisionTable

163249 - rules.JRip

Classifier output

Time taken to build model: 0.03 seconds

==== Stratified cross-validation ====

==== Summary ====

	Correctly Classified Instances	53	92.9825 %
	Incorrectly Classified Instances <td>4 <td>7.0175 % </td></td>	4 <td>7.0175 % </td>	7.0175 %
	Kappa statistic <td>0.8494</td> <td></td>	0.8494	
	Mean absolute error <td>0.0641</td> <td></td>	0.0641	
	Root mean squared error <td>0.2438</td> <td></td>	0.2438	
	Relative absolute error <td>14.0069 %</td> <td></td>	14.0069 %	
	Root relative squared error <td>51.0652 %</td> <td></td>	51.0652 %	
	Total Number of Instances <td>57</td> <td></td>	57	

==== Detailed Accuracy By Class ====

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.950	0.081	0.864	0.950	0.905	0.852	0.970	0.927	bad
	0.919	0.050	0.971	0.919	0.944	0.852	0.981	0.989	good
Weighted Avg.	0.930	0.061	0.934	0.930	0.931	0.852	0.977	0.967	

==== Confusion Matrix ====

	a	b	<-- classified as
19	1	1	a = bad
3	34	1	b = good

Status

OK

Log

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier Choose JRip - F 3 - N 2.0 - O 2 - S 1

Test options

☐ Use training set

☐ Supplied test set

☒ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 135149 - bayes.NaiveBayes
- 135220 - bayes.NaiveBayesMultinomialText
- 135425 - bayes.NaiveBayesUpdateable
- 162623 - functions.Logistic
- 162645 - functions.MultilayerPerceptron
- 162746 - meta.AdaBoostM1
- 162808 - meta.Bagging
- 163008 - trees.RandomTree
- 163044 - trees.J48
- 163228 - rules.DecisionTable
- 163249 - rules.JRip

Classifier output

Attrib contribution to health-plan=full -0.19217350105123066

Class bad

Input

Node 0

Class good

Input

Node 1

Time taken to build model: 0.23 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances	49	85.9649 %
Incorrectly Classified Instances	8	14.0351 %
Kappa statistic	0.6919	
Mean absolute error	0.1524	
Root mean squared error	0.337	
Relative absolute error	33.3246 %	
Root relative squared error	70.5761 %	
Total Number of Instances	57	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.800	0.108	0.800	0.800	0.800	0.692	0.923	0.917	bad
	0.892	0.200	0.892	0.892	0.892	0.692	0.923	0.951	good
Weighted Avg.	0.860	0.168	0.860	0.860	0.860	0.692	0.923	0.939	

=== Confusion Matrix ===

a b <-- classified as

16	4	a = bad
4	23	b = good

Status OK

Log

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier Choose JRip - F 3 - N 2.0 - O 2 - S 1

Test options

☐ Use training set

☐ Supplied test set

☒ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 135149 - bayes.NaiveBayes
- 135220 - bayes.NaiveBayesMultinomialText
- 135425 - bayes.NaiveBayesUpdateable
- 162623 - functions.Logistic
- 162645 - functions.MultilayerPerceptron
- 162746 - meta.AdaBoostM1
- 162808 - meta.Bagging
- 163008 - trees.RandomTree
- 163044 - trees.J48
- 163228 - rules.DecisionTable
- 163249 - rules.JRip

Classifier output

ways instance class year is missing

bad good

0.0 1.0

Weight: 1.12

Number of performed iterations: 10

Time taken to build model: 0.01 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances	50	87.7193 %
Incorrectly Classified Instances	7	12.2807 %
Kappa statistic	0.7273	
Mean absolute error	0.1519	
Root mean squared error	0.3414	
Relative absolute error	33.2028 %	
Root relative squared error	71.5013 %	
Total Number of Instances	57	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.800	0.081	0.842	0.800	0.821	0.728	0.870	0.863	bad
	0.919	0.200	0.855	0.919	0.907	0.728	0.870	0.876	good
Weighted Avg.	0.877	0.158	0.876	0.877	0.876	0.728	0.870	0.871	

=== Confusion Matrix ===

a b <-- classified as

16	4	a = bad
3	34	b = good

Status OK

Log

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose **JRip - F 3 -N 2.0 -O 2 -S 1**

Test options

☐ Use training set

☐ Supplied test set

☒ Cross-validation Folds **10**

☐ Percentage split % **66**

More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 135149 - bayes.NaiveBayes
- 135220 - bayes.NaiveBayesMultinomialText
- 135425 - bayes.NaiveBayesUpdateable
- 162623 - functions.Logistic
- 162645 - functions.MultilayerPerceptron
- 162746 - meta.AdaBoostM1
- 162808 - meta.Bagging**
- 163008 - trees.RandomTree
- 163044 - trees.J48
- 163228 - rules.DecisionTable
- 163249 - rules.JRip

Classifier output

Test mode: 10-fold cross-validation

==== Classifier model (full training set) ====

Bagging with 10 iterations and base learner

weka.classifiers.trees.REPTree -M 2 -V 0.001 -M 3 -S 1 -L -I 0.0

Time taken to build model: 0.03 seconds

==== Stratified cross-validation ====

==== Summary ====

Correctly Classified Instances	48	84.2105 %
Incorrectly Classified Instances	9	15.7895 %
Kappa statistic	0.641	
Mean absolute error	0.2979	
Root mean squared error	0.3757	
Relative absolute error	65.1224 %	
Root relative squared error	78.6822 %	
Total Number of Instances	57	

==== Detailed Accuracy By Class ====

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.700	0.081	0.824	0.700	0.757	0.646	0.856	0.774	bad
	0.919	0.300	0.850	0.919	0.883	0.646	0.856	0.886	good
Weighted Avg.	0.842	0.223	0.841	0.842	0.839	0.646	0.856	0.847	

==== Confusion Matrix ====

a	b	<-- classified as
14	6	a = bad
3	34	b = good

Status OK

Log

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier

Choose **JRip - F 3 -N 2.0 -O 2 -S 1**

Test options

☐ Use training set

☐ Supplied test set

☒ Cross-validation Folds **10**

☐ Percentage split % **66**

More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 135149 - bayes.NaiveBayes
- 135220 - bayes.NaiveBayesMultinomialText
- 135425 - bayes.NaiveBayesUpdateable
- 162623 - functions.Logistic
- 162645 - functions.MultilayerPerceptron
- 162746 - meta.AdaBoostM1
- 162808 - meta.Bagging
- 163008 - trees.RandomTree**
- 163044 - trees.J48
- 163228 - rules.DecisionTable
- 163249 - rules.JRip

Classifier output

Test mode: 10-fold cross-validation

==== Classifier model (full training set) ====

Size of the tree : 30

Time taken to build model: 0.01 seconds

==== Stratified cross-validation ====

==== Summary ====

Correctly Classified Instances	45	78.9474 %
Incorrectly Classified Instances	12	21.0526 %
Kappa statistic	0.5156	
Mean absolute error	0.2438	
Root mean squared error	0.4353	
Relative absolute error	53.2875 %	
Root relative squared error	91.1643 %	
Total Number of Instances	57	

==== Detailed Accuracy By Class ====

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.600	0.108	0.750	0.600	0.667	0.522	0.754	0.654	bad
	0.892	0.400	0.805	0.892	0.846	0.522	0.754	0.795	good
Weighted Avg.	0.789	0.298	0.786	0.789	0.783	0.522	0.754	0.746	

==== Confusion Matrix ====

a	b	<-- classified as
12	8	a = bad
4	33	b = good

Status OK

Log

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier Choose J48 - F 3 - N 2.0 - O 2 - S 1

Test options

- ☐ Use training set
- ☐ Supplied test set
- ☒ Cross-validation Folds 10
- ☐ Percentage split % 66

More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 135149 - bayes.NaiveBayes
- 135220 - bayes.NaiveBayesMultinomialText
- 135425 - bayes.NaiveBayesUpdateable
- 162623 - functions.Logistic
- 162645 - functions.MultilayerPerceptron
- 162746 - meta.AdaBoostM1
- 162808 - meta.Bagging
- 163008 - trees.RandomTree
- 163044 - trees.J48
- 163228 - rules.DecisionTable
- 163249 - rules.JRip

Classifier output

```
wage-increase-first-year <= 2.0: bad (10.77/4.77)
wage-increase-first-year > 2.0
| statutory-holidays <= 10: bad (10.77/4.77)
| statutory-holidays > 10: good (30.96/1.0)
```

Number of Leaves : 3
Size of the tree : 5

Time taken to build model: 0.02 seconds

=== Stratified cross-validation ===
=== Summary ===

Correctly Classified Instances	42	73.6842 %
Incorrectly Classified Instances	15	26.3158 %
Kappa statistic	0.4415	
Mean absolute error	0.3192	
Root mean squared error	0.4669	
Relative absolute error	69.7715 %	
Root relative squared error	97.7888 %	
Total Number of Instances	57	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.700	0.243	0.609	0.700	0.651	0.444	0.695	0.559	bad
	0.757	0.300	0.824	0.757	0.789	0.444	0.695	0.738	good
Weighted Avg.	0.737	0.280	0.748	0.737	0.740	0.444	0.695	0.675	

=== Confusion Matrix ===

a	b	<-- classified as
14	6	a = bad
9	28	b = good

Status OK

Log

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier Choose J48 - F 3 - N 2.0 - O 2 - S 1

Test options

- ☐ Use training set
- ☐ Supplied test set
- ☒ Cross-validation Folds 10
- ☐ Percentage split % 66

More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 135149 - bayes.NaiveBayes
- 135220 - bayes.NaiveBayesMultinomialText
- 135425 - bayes.NaiveBayesUpdateable
- 162623 - functions.Logistic
- 162645 - functions.MultilayerPerceptron
- 162746 - meta.AdaBoostM1
- 162808 - meta.Bagging
- 163008 - trees.RandomTree
- 163044 - trees.J48
- 163228 - rules.DecisionTable
- 163249 - rules.JRip

Classifier output

```
Start set: no attributes
Search direction: forward
Stale search after 5 node expansions
Total number of subsets evaluated: 84
Merit of best subset found: 92.982
Evaluation (for feature selection): CV (leave one out)
Feature set: 2,13,17
```

Time taken to build model: 0.05 seconds

=== Stratified cross-validation ===
=== Summary ===

Correctly Classified Instances	43	75.4386 %
Incorrectly Classified Instances	14	24.5614 %
Kappa statistic	0.4729	
Mean absolute error	0.2996	
Root mean squared error	0.4014	
Relative absolute error	65.5028 %	
Root relative squared error	84.0671 %	
Total Number of Instances	57	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.700	0.216	0.636	0.700	0.667	0.474	0.818	0.722	bad
	0.784	0.300	0.829	0.784	0.806	0.474	0.818	0.884	good
Weighted Avg.	0.754	0.271	0.761	0.754	0.757	0.474	0.818	0.827	

=== Confusion Matrix ===

a	b	<-- classified as
14	6	a = bad
8	29	b = good

Status OK

Log

Weka Explorer

PreprocessClassifyClusterAssociateSelect attributesVisualize

Classifier

ChooseJRip-F 3-N 2.0-O 2-S 1

Test options

☐ Use training set

☐ Supplied test set

☒ Cross-validation

☐ Percentage split

Folds10

%66

More options...

(Nom) class

Start

Stop

Result list (right-click for options)

13:51:49 - bayes.NaiveBayes

13:52:20 - bayes.NaiveBayesMultinomialText

13:54:25 - bayes.NaiveBayesUpdateable

16:26:23 - functions.Logistic

16:26:45 - functions.MultilayerPerceptron

16:27:46 - meta.AdaBoostM1

16:28:08 - meta.Bagging

16:30:08 - trees.RandomTree

16:30:44 - trees.J48

16:32:28 - rules.DecisionTable

16:32:49 - rules.JRip

Classifier output

(wage-increase-first-year <= 2.5) => class=bad (15.0/2.0)

(statutory-holidays <= 10) and (wage-increase-first-year <= 4) => class=bad (5.0/0.0)

(longterm-disability-assistance = no) => class=bad (2.0/0.0)

=> class=good (35.0/0.0)

Number of Rules : 4

Time taken to build model: 0.03 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances4477.193 %

Incorrectly Classified Instances1322.807 %

Kappa statistic0.4935

Mean absolute error0.2256

Root mean squared error0.456

Relative absolute error49.3094 %

Root relative squared error95.4949 %

Total Number of Instances57

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
	0.650	0.162	0.684	0.650	0.667	0.494	0.779	0.718	bad
	0.838	0.350	0.816	0.838	0.827	0.494	0.779	0.810	good
Weighted Avg.	0.772	0.284	0.770	0.772	0.771	0.494	0.779	0.778	

=== Confusion Matrix ===

a b

<-- classified as

13 7 | a = bad

6 31 | b = good

Status

OK

Log

x 0

Windows Search

Taskbar icons: File Explorer, Microsoft Edge, Google Chrome, VS Code, Outlook, Word, PowerPoint, Teams, OneDrive, etc.

System tray: Network, Volume, ENG IN, Date/Time: 16:51 13-06-2024